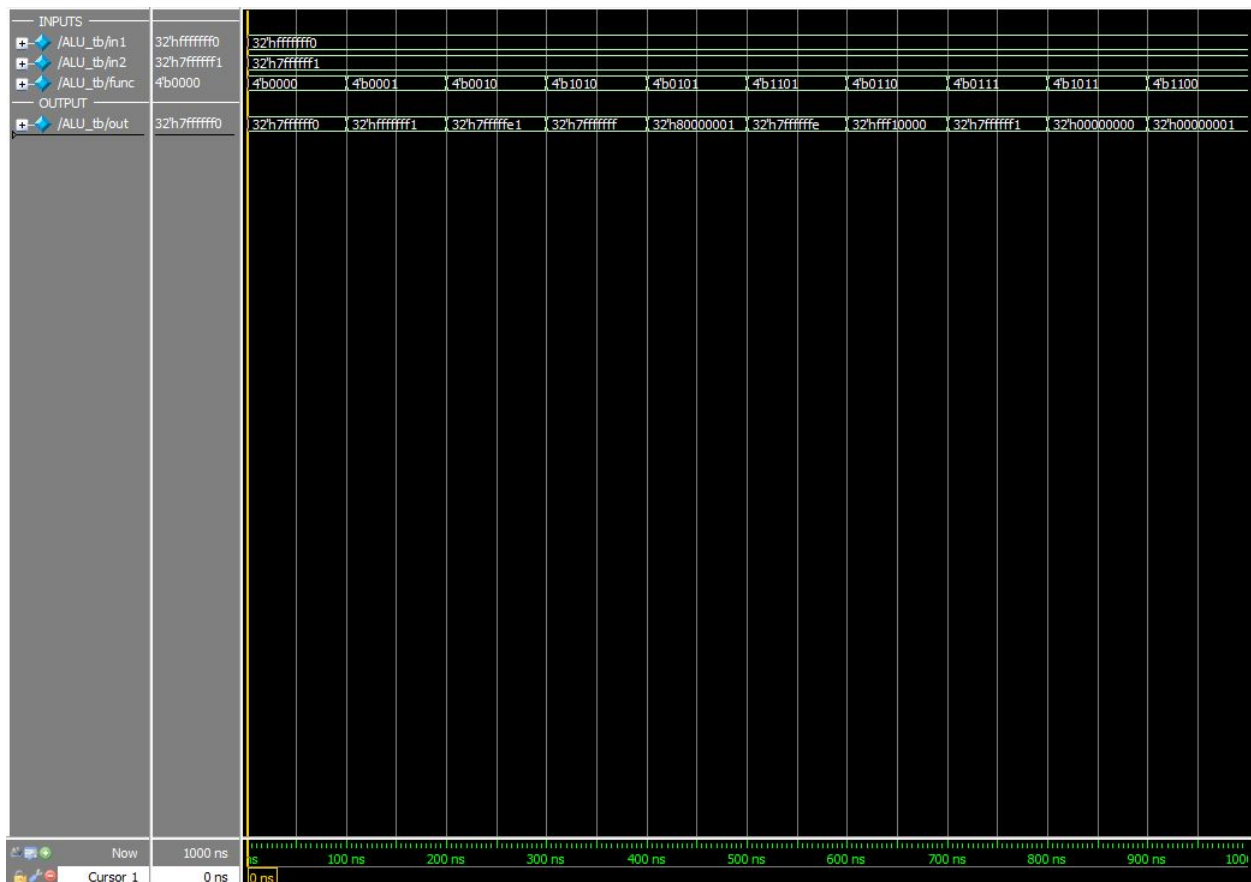


Introduction

The idea of a pipeline processor is to utilize more parts of the processor at the same time. The full operation can be split into multiple stages such that the same operations done multiple times can be overlapped together and executed simultaneously. This design is broken into 5 stages: fetch, decode, execute, memory, and writeback. One of the main problems with pipelined registers is that dependent instructions that are not finished executing results in hazards. The hazard unit detects and resolves the hazards.

ALU


The ALU's basic operations are add, sub, and, or, slt, lui, xor and in2 for mfhi and mflo. The last bit in ALUop is designed to negate in2 most notably for sub and xor. Unsigned operations are the same representation wise except for slt. An extra operation configuration is used for sltu to check the sign bits.



ALU_tb waveform

Data Memory

The data memory holds 64 slots of 32 bit words. Each word is byte addressable, but indexing in Verilog is every integer, so the address is shifted to the right twice to align with byte addressing. Writing to the memory is on the positive edge and the enable signal needs to be on.



data_memory_tb waveform

Instruction Memory


The instruction memory can hold 64 32 bit instructions. Addressing into the instruction memory is the same as data memory. Initializing instructions into the memory uses memfile.dat. The figure shows the first 10 memory slots initialized with constants from 0-9.

[illegible]

Instruction Memory

Regfile


The register file contains 32 32 bit registers. On reset, all the registers are reset to 0. Writing on the positive edge requires an enable signal. Two read ports are asynchronous and should update if data is written into that register in the same cycle.



reg_file_tb

Controller


The controller outputs all the enable signals, MUX selects, and ALUOp. Op is used first, then funct if the instruction is an R type. Eq_ne is from the equality check for early branch prediction. The default output for all 0s instructions are essentially nops.

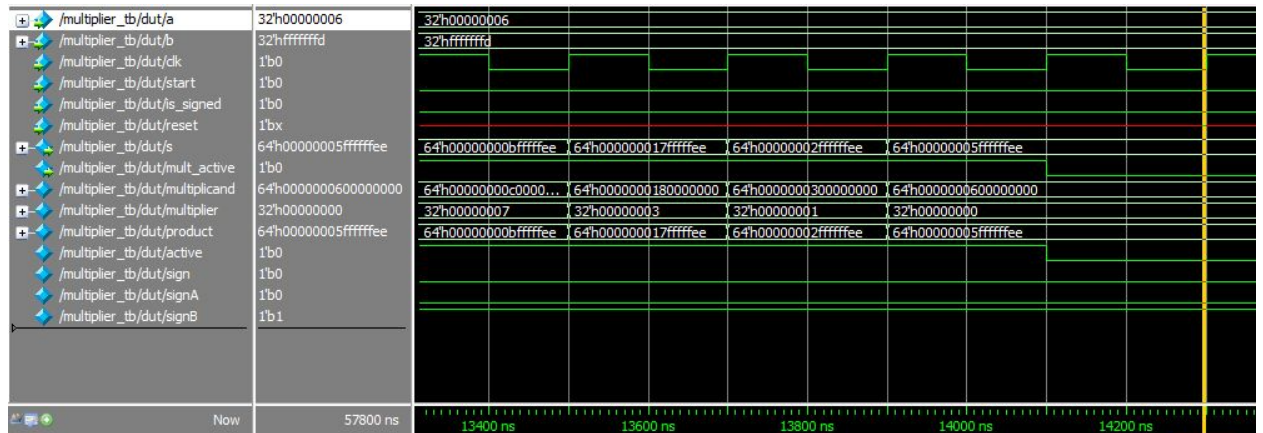



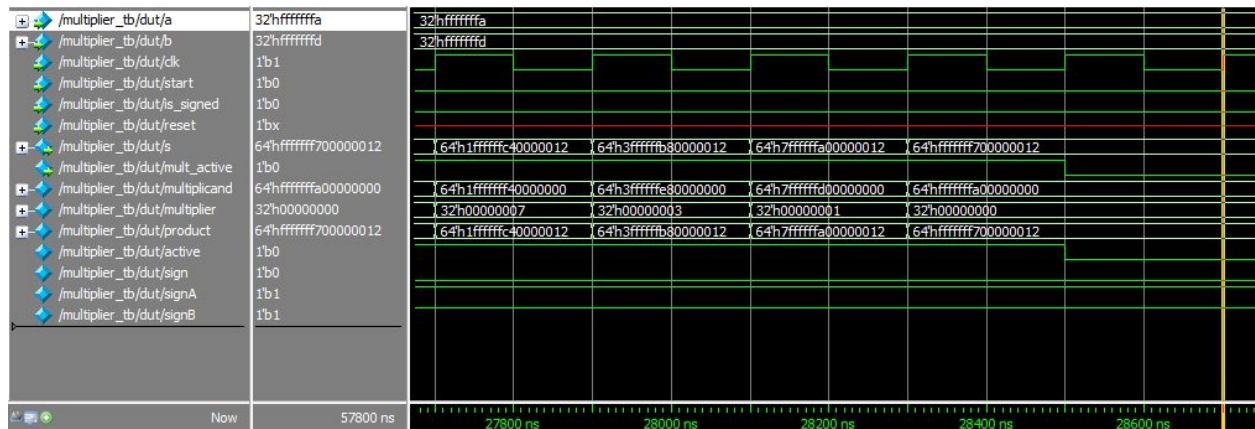
controller_tb waveform

Multiplier

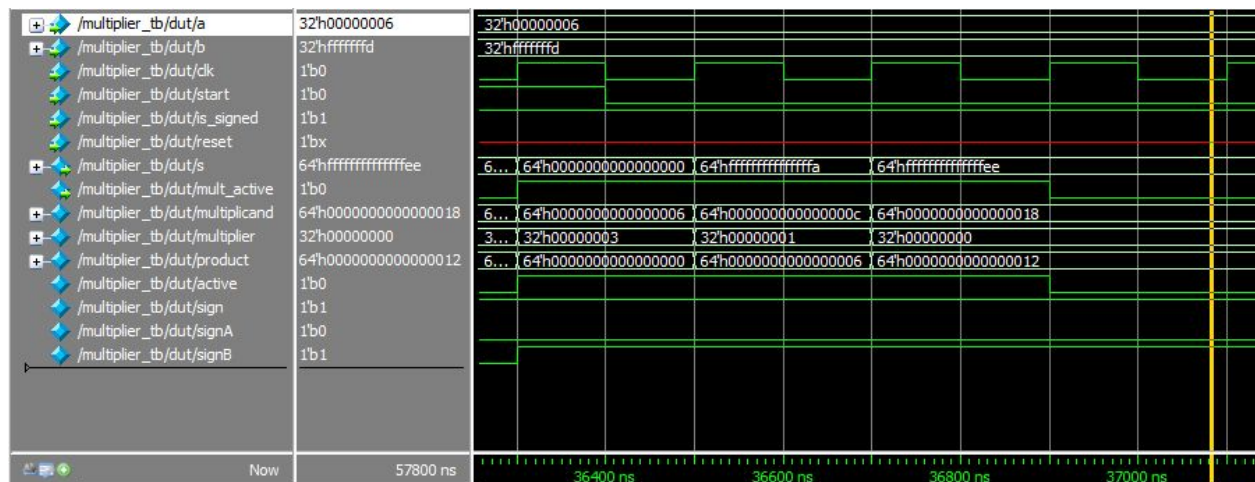
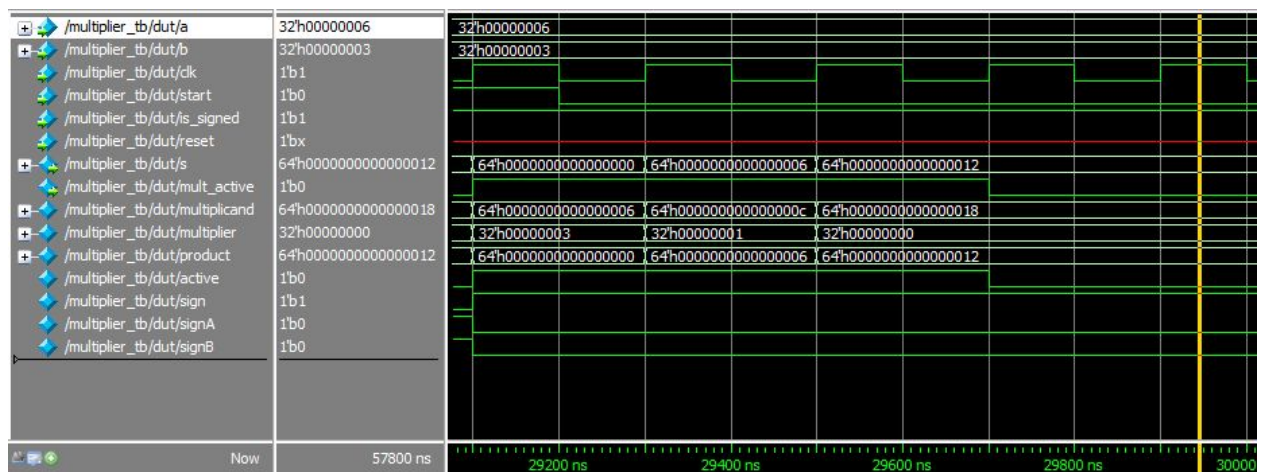
A serial multiplier takes the two inputs and adds up their partial sums. The multiplicand is shifted to the left and added into the product. The multiplier is shifted to the right and the LSB is checked if the addition is needed. For signed multiplication, the unsigned multiplication is performed first, then the sign is added if needed. An active signal is outputted for the processor to stall until the multiplication is finished. The serial multiplier finishes as soon as the multiplier is shifted to all 0s. The product can be 64 bits long, so it is stored in its own location. To move the value to the regfile, mfhi and mflo will move the high 32 bits and low 32 bits respectively. The first four figures are unsigned and the last four are signed.

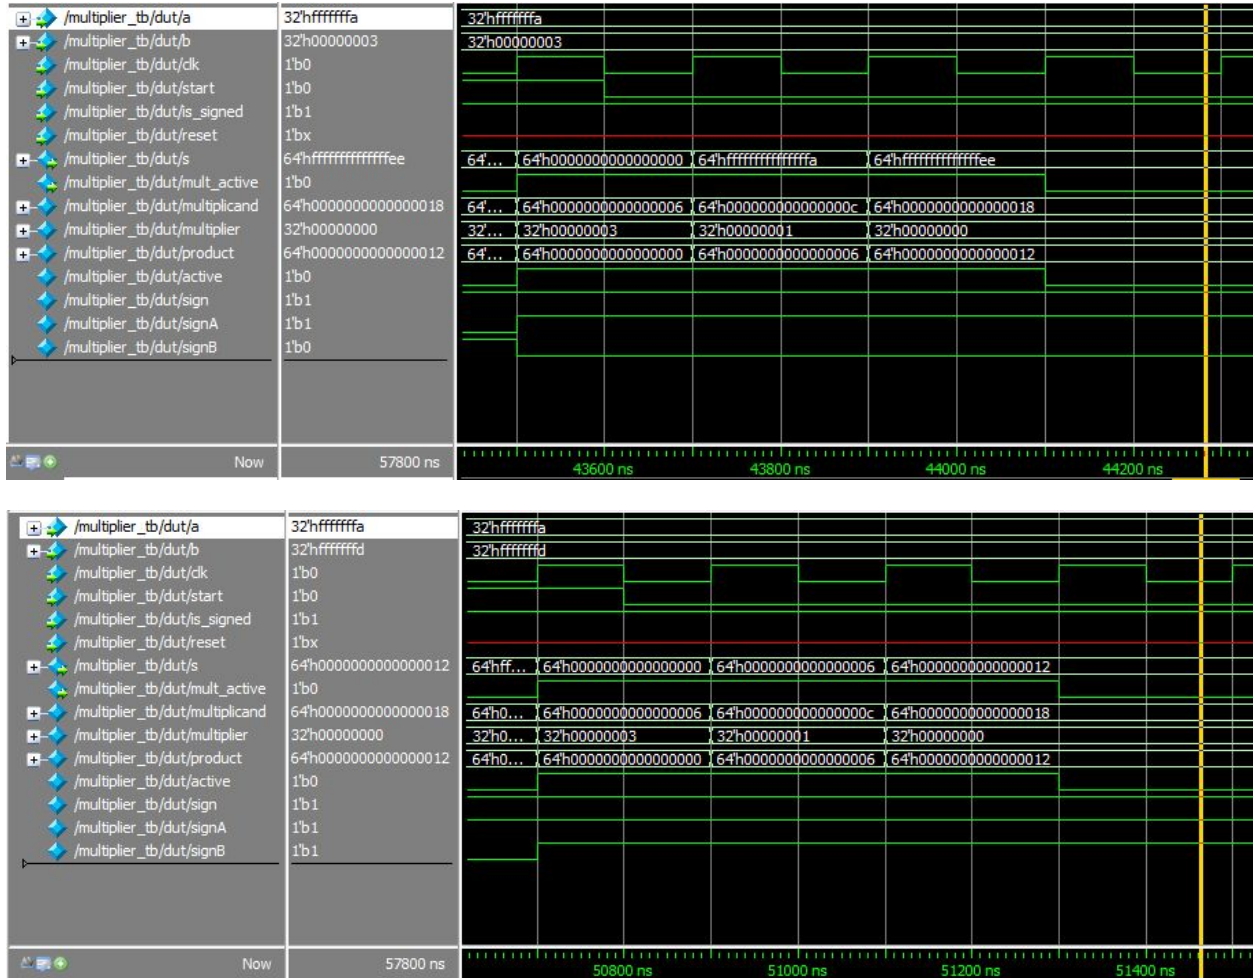






multu waveforms







mult waveform

Hazard Unit

If the subsequent instruction is dependent on the previous registers, but the value has not been written back yet (stage 5), then the data needs to be forwarded to the correct stage. If a value is needed from memory right after a lw instruction, the processor is forced to stall until the data is read from memory. With early branch prediction, the processor knows when to take the branch at the decode stage, but this causes another data hazard. The processor needs to stall, then forward the data to the decode stage if needed. The fetch and decode registers are stalled and execute registers are flushed to prevent unwanted instructions from being executed. The first figure shows data forwarding and the second shows all types of stalling.



hazard_unit_tb waveform

Datapath

Some notes on datapath:

ALU Src has 2 bits to take into account `mfhi` and `mflo` as inputs. Connecting writing ports of regfile to writeback registers instead of the wires at the end of the memory stage delays the writing by one cycle. The controller in theory has don't care terms as output, but it is safer to set a known value to those control signals because it messes up values for future cycles like stall signals for example. On reset, all the register files are set to all 0s. This is essentially a nop.

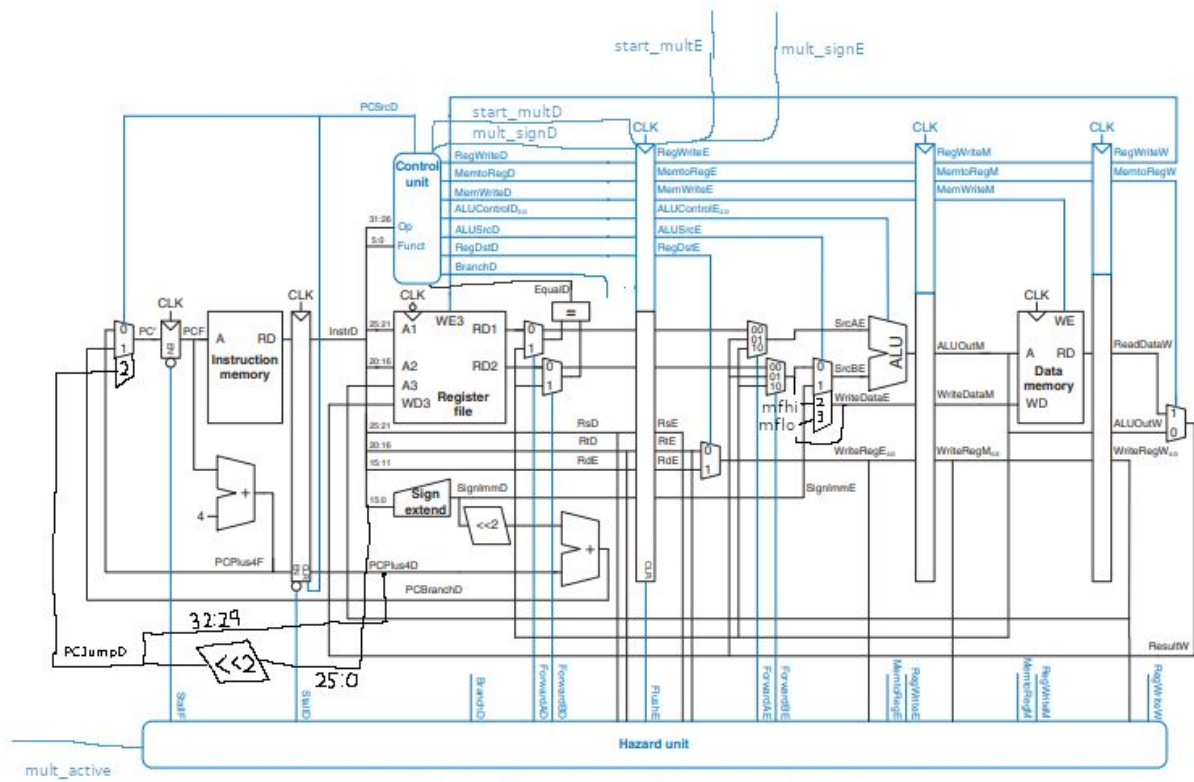


Figure 7.58 Pipelined processor with full hazard handling

Memfile1.dat

This program is the same as the one in figure 7.60 on page 437. The regfile and memory should have these values:

\$2 = 7

\$3 = 12

\$4 = 1

\$5 = 11

\$7 = 7

mem[80] = \$7 = 7

mem[84] = \$2 = 7

0000003f	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
00000037	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
0000002f	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
00000027	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
0000001f	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
00000017	XXXXXXXX	XXXXXXXX	00000007	00000007	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
0000000f	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
00000007	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX


Data memory


0000001f	00000000	00000000	00000000	00000000	00000000	00000000	00000000	00000000
00000017	00000000	00000000	00000000	00000000	00000000	00000000	00000000	00000000
0000000f	00000000	00000000	00000000	00000000	00000000	00000000	00000000	00000000
00000007	00000003	00000000	0000000b	00000007	0000000c	00000005	00000000	00000000

Regfile

Memfile2.dat

This program finds the dot product of two vectors. The vectors are not initialized in memory, but one element from each vector is stored at a time in the regfile. The answer is 0d4985 or 0h1379. It takes 75 cycles to write the dot product back into the regfile.





Waveform for computing product, then summation of second element.

Memfile3.dat

For two arrays in memory of n length, $Y[j+1] = B \cdot X[j] + Y[j]$. Instead of using the multiplier for $B \cdot X[j]$, the program instead adds $X[j]$ B times as a running sum. B and n are initialized in $\$a0$ and $\$a1$ respectively. The first 10 memory slots are for X and the next 11 are for Y .

0000003f	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
00000037	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
0000002f	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
00000027	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
0000001f	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX	XXXXXXXX
00000017	XXXXXXXX	XXXXXXXX	XXXXXXXX	00000052	0000004d	0000003e	0000001b	0000001b	
0000000f	0000003e	00000020	00000011	0000001b	fffffffe	00000002	00000001	00000003	
00000007	00000007	00000000	fffffff9	00000006	00000003	fffffffe	00000009	fffffffc	

Data Memory

Conclusion

Pipeline processors are faster than single cycle and multi cycle processors because they can execute multiple instructions at a time. The limiting factors are the amount of hazards as the number of stages increase and the cost of setup and hold times between every pipeline register. Increasing the number of stages might slow down the processor due to the increase in hazards

and pipeline register overhead. All the hazards force the processor to stall for a number of cycles and the goal is to decrease the number of stalls as much as possible.